

Data Science Career Track

Model Metrics Exercise

1. Look at the table below. If the goal is to optimize the True Positives, which model would you choose and why?

Model	Recall	Precision	Accuracy	F1
Logistic	0.746	0.775	0.999	0.761
Logistic with auto threshold	0.891	0.061	0.976	0.114
Logistic with class weights	0.878	0.110	0.988	0.195
Hinge with auto threshold	0.905	0.014	0.890	0.028
Hinge with class weights	0.878	0.103	0.987	0.185

- a. If we want to optimize the True Positives, we should choose the model with the highest Precision relative to other models. Here, that would be the default Logistic Regression model, with a Precision score of 0.775, Accuracy of 0.999, and F1-score of 0.761.
2. Calculate the F-1 scores for each model and identify the best model based on the F1 score.

Model	Recall	Precision	F1	Auc/Roc
Deep NN	0.79	0.82	0.80	0.92
Logistic Regression	0.75	0.79	0.77	0.90
Random Forest	0.80	0.66	0.72	0.90
LinearSVC	0.74	0.75	0.74	0.82

- a. Calculated using the following:

$$\text{F1 score} = \frac{2 * \text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

3. Identify the best parameter values for 'alpha' and 'L1-ratio' based on the above comparison.

Model	Parameter	Parameter	Metric	Metric	Metric
	Alpha	L1-ratio	MAE	R-squared	RMSE
Linear Regression	0.5	0.2	84.27	0.277	158.1
Linear Regression	0.2	0.5	84.08	0.264	159.6
Linear Regression	0.5	0.5	84.12	0.272	158.6
Linear Regression	0	0	84.49	0.249	161.2

- a. Alpha = 0.5, L1-ratio = 0.2.

- i. While these parameters do not produce the best (lowest) MAE across the iterations, they *do* yield the best (highest) R-squared and the best (lowest) RMSE.